

Workshop Report: Creating a Template TVEC Policy for Text Classification


christchurchcall.org

ELEVATE

23 April 2026

Table of Contents

Executive Summary	3
Introduction	4
Feedback Collected	5
Policy Revisions	7
Future Areas of Work: CTVE101 Toolkit	8



Executive Summary

In October 2025, ROOST and OpenAI partnered to open-source gpt-oss-safeguard, a reasoning model fine-tuned for safety applications. Gpt-oss-safeguard is unique because of its “bring-your-own policy” approach, allowing those who use the tool to set their own definitions of allowable and violative content.

Following the release of this model, the Christchurch Call Foundation (CCF) began working on a “template” policy for terrorist and violent extremist content (TVEC) that would support effective application of gpt-oss-safeguard and similar safety reasoning models for text classification. After creating a first draft of this policy in December 2025, the CCF team worked with members of the Christchurch Call’s multistakeholder community to test the policy in a sandbox environment and gather recommendations for improvement.

On 21 January 2026, the Christchurch Call Foundation held two workshops to collect feedback from our multistakeholder community on the development of this policy. These live feedback sessions were complemented by a written feedback form. The purpose of this report is to share a summary of the feedback we collected and how it informed improvements to the policy.

Key Findings

Together, the feedback collected from the two workshop sessions and written form surfaced the following five key findings and recommendations:

1. The policy could be strengthened by including a list of terrorist and violent extremist actors to mitigate errors stemming from model knowledge gaps.
2. In-workshop testing surfaced a need to refine the exceptions categories, which the model often applied inaccurately and were easily gamed by participants.
3. Including severity scores and escalation recommendations within the policy would help trust and safety teams prioritize enforcement action and prevent harm.
4. Integrate more precise definitions and examples of terrorist content and terrorist activity online.
5. In-workshop testing indicated inconsistency in model outputs in languages other than English; an issue that is more dependent on model training rather than policy fine-tuning.

Introduction

Workshop participants were asked to keep the following considerations in mind as they reviewed and tested the draft policy:

- **The template TVEC policy is designed specifically to be effective when used alongside open reasoning models for text classification.** This reasoning model is not intended to be a “one-stop-shop” for content moderation. Effectively protecting online communities from TVEC necessitates tools that support detection, investigation, review, and enforcement. This tool and our accompanying TVEC policy are just one piece of the puzzle.
- **“Bring-your-own policy” reasoning models like gpt-oss-safeguard work best when provided with clear definitions and examples.** Drafting these examples and definitions was challenging – as there are no universally agreed definitions of TVE and TVEC. However, failure to provide reasoning models with clear definitions may introduce ambiguity that results in the classification of protected speech as TVEC. We believe the best approach is to use definitions that have been informed by expertise from across the Call’s multistakeholder community.
- **The template TVEC policy is first and foremost an operational policy designed for content moderation and enforcement.** It is meant to support the use of open-source trust and safety tools – specifically gpt-oss-safeguard – for identifying TVEC text content. As such, this policy serves a different purpose from public facing TVE policies often found in online service providers’ Community Guidelines which explain their approach to countering TVE on their services and set expectations for community behaviour.
- **The template TVEC policy is just that – a template.** It should offer a strong blueprint for users of open-source trust and safety tools while also offering adaptability to diverse services, threat types, cultural and political contexts, and approaches to moderation. In that sense, the policy is intended to be a practical base foundation to support tech companies’ own trust and safety workflows, rather than a prescriptive approach to countering terrorism and violent extremism online.

Importantly, this policy is designed specifically to support classification of text content. In the coming year, the Elevate team will be working to expand on this effort by developing a catalogue of TVE policies tailored to different open-source trust and safety functions (i.e., multi-modal content, behavioural signals, network investigations) and content moderation contexts (i.e., entertainment, creative, and social content).

Method for Community Engagement

After drafting an initial version of the template TVEC policy, CCF engaged experts in the Call Community to collect feedback on areas for improvement. The CCF team hosted two workshop sessions on 21 January 2026. Together, these were attended by 23 experts across government, civil society, and industry. Community members also had the opportunity to provide written feedback via a Microsoft form. We received feedback from one respondent via the written form.

Feedback Collected

At the beginning of the workshop, participants were asked what values they thought should guide the development of TVEC policies for open-source safety reasoning models like gpt-oss-safeguard. The answers collected provide a snapshot of the diverse perspectives within the Christchurch Call's multistakeholder community. These values provide guidance to the CCF team as we seek to update and maintain TVEC policy resources over time. The values identified through this exercise can be grouped into five core themes:

- Proportionality, transparency, and fairness
- Scalability and adaptability
- Respect for international law, human rights, and freedom of expression
- Clarity
- Cultural sensitivity

Following the identification of guiding values, participants were provided with an overview of the policy and invited to spend time working in small groups to test the draft policy on example pieces of text in a sandbox environment on Hugging Face. After spending 20 minutes testing the policy in small groups, CCF facilitated a discussion where participants were asked to share their findings and recommendations for strengthening the draft policy. The key takeaways from these discussions have been combined with feedback collected via the written form and summarised into the five key findings below.

Finding 1: The policy could be strengthened by including a list of terrorist and violent extremist actors to mitigate errors stemming from model knowledge gaps.

Participants in the workshops found that the model was inconsistent when classifying content referencing terrorist actors or groups. While the model appeared to have a more robust understanding of violent extremism in the US context, for example, it was inconsistent in recognising violent extremist actors in other contexts, including Kenya and Ireland. The model was also inconsistent in recognising terrorist actors when acronyms or abbreviations were used. As a result, participants recommended that the policy be paired with a list of TVE actors that teams applying gpt-oss-safeguard for CTVE could reference. Importantly, however, it was noted that reference to a specific TVE organization or individual perpetrator should be clearly articulated as one signal among many, rather than considered a policy violation on its own.

Participants were often able to circumvent the policy by using slang or common phrases associated with violent extremist ideologies. While policies are not the right place for such datasets, these findings suggest that the application of this tool could be strengthened by a complementary signals dataset that draws on existing resources across the field of TVE research. Participants noted that signals are needed to respond to changing nature of traditional TVE threats as well as emerging hybrid harms.

Finding 2: In-workshop testing surfaced a need to refine the exceptions categories, which the model often applied inaccurately and were easily gamed by participants.

ELEVATE

Challenges with the exemption categories were found by participants in both workshops. For example, when clearly violative statements were falsely claimed to be written by a reputable news outlet, the gpt-oss-safeguard model often returned a TVEC exemption label.

Recommendations for strengthening these exemptions include tightening the definition for “newsworthiness” and shifting the exceptions from relying on the format and author to the presentation of the content. For example, the way that TVEC.02 – Journalistic Content Exemption was originally written relied too much on the model being able to discern who is and is not a reputable journalist and introduced opportunity for gaming.

While some of these gaps can be addressed through updates to the template policy, it is important to note that content alone can often be insufficient when classifying terrorism and violent extremism online. Often, content needs to be considered alongside actor and behavioural signals to better understand the context around which content is shared and make an accurate assessment of whether it constitutes terrorism and violent extremism. The policy for gpt-oss-safeguard deals with the content piece of this puzzle, specifically, and we note that it is up to the discretion of online service providers to apply this policy in context with other signals.

Finding 3: Including severity scores and escalation recommendations within the policy would help trust and safety teams prioritize enforcement action and prevent harm.

Participants felt that the recommended prescriptions of the policy (allowed versus not allowed) were too broad, and that the template policy could go a step further by including an indication of severity and recommended actions beyond removal (i.e., reporting to law enforcement). This was originally considered by the CCF team but was ultimately left out of the draft policy given the variation in reporting requirements across jurisdictions.

However, there are ways that we could add some indication of severity while avoiding tailoring the policy too closely to legal jurisdictions. As a first step, participants suggested that this could be achieved by providing a severity rating for each category that is based on likelihood of violence.

Finding 4: Integrate more precise definitions and examples of terrorist content and terrorist activity online.

Participants noticed a tension in the scope of the policy as the definitions used for some categories were highly tailored towards content, whereas others included examples and indicators of related more to terrorist or violent extremist activity. The current policy is intended to focus only on text content, to maximize its suitability for gpt-oss-safeguard. This scope ultimately limits the policy’s applicability to harm areas like material support, recruitment, promotion and incitement. Designing policies that respond to content, actor, and behavioural signals is particularly important when responding to risks introduced by hybridized threats, including nihilistic violent extremism (NVE) and school massacre ideation.

Finding 5: In-workshop testing indicated inconsistency in model outputs in languages other than English; an issue that is more dependent on model training rather than policy fine-tuning.

Participants found that the model sometimes produced a different assessment of the same statement written when tested in English versus other languages (languages tested during the workshop included French, Irish, Dutch, Swahili, and more). In one specific example, a statement

ELEVATE

glorifying ISIS in Swahili was not labelled as a policy violation, whereas the same sentence presented in English was detected as a policy violation. Further, the same statement, written in Swahili but glorifying Boko Haram was detected as a policy violation.

This reveals inconsistency in the model's capacity to understand TVEC across linguistic contexts and is representative of language limitations across online safety tools. However, this limitation is a product of model training, rather than the configuration of our TVEC policy.

Policy Revisions

The following changes were made to the policy to improve performance following consultation with experts from our multistakeholder community:

- To support a more effective application of the policy, particularly by teams who do not have dedicated subject matter expertise in terrorism and violent extremism, we have included an Annex which provides a reference list of designated and proscribed terrorist organizations as well as the names of individual perpetrators of recent violent extremist attacks. This Appendix refers to designated and proscribed lists from the United Nations, the European Union, Australia, Canada, New Zealand, the United Kingdom, and the United States. We recognize that there is no single, universally agreed list of designated and proscribed entities. However, we hope that this can be a starting point for teams with limited knowledge of the existing field, upon which considerations for regional contextualization can be expanded. For the list of individual perpetrators, we used a list based on key terrorist and violent extremist incidents where 1) the Christchurch Call Foundation engaged its Crisis Response Protocol or 2) the incident involved references to the Christchurch attack. As such, this list is not all-encompassing, but rather serves as an initial starting point.
- The revised policy includes significant updates to the TVEC exceptions to reduce the potential for gaming, responding to flaws identified by workshop participants. This includes a more explicit articulation of when exceptions do not apply. Exceptions have been condensed from three categories (EDSA, journalistic content, and newsworthiness) to one category. This change was made because testing revealed that the model often struggled to distinguish between the three different exceptions categories.
- The revised policy includes a clear threat level indicator in the label produced by the model to indicate to human reviewers the severity of the content.
- The policy includes refined examples and categories that clearly match the application of the policy for terrorist and violent extremist text content only. This scope is also clearly indicated in how the policy is presented on GitHub. CCF received excellent feedback from participants about the need to develop policies that respond to multimodal content and consider actor, network-level, and behavioural signals. These areas will be addressed with the development of additional policies within our forthcoming CTVE101 Toolkit, which is described in greater detail in the Future Areas of Work section below.
- The documentation on GitHub supporting the use of the policy includes a note describing issues when using gpt-oss-safeguard in languages other than English. Inconsistent results produced across languages can only be solved through model training, rather than policy.

Future Areas of Work: CTVE101 Toolkit

In addition to the revisions described above, the workshop feedback surfaced the need for additional resources to support effective and robust counter-terrorism and violent extremism (CTVE) operations online.

To support the deployment of open safety solutions for CTVE operations, the Christchurch Call Foundation will develop a CTVE 101 Toolkit comprising of:

- **A repository of open source and accessible trust and safety solutions suited for CTVE operations**, organised across building blocks from minimum viable content detection to full CTVE suite.
- **A CTVE policy catalogue**, comprising of practical and baseline policies adapted to different online safety tools and trust and safety functions, from multimodal content detection to actor level and platform wide investigation.
- **A CTVE handbook**, including key information on how terrorists and violent extremists use online services and implementation guidance to support the deployment of the open safety solutions listed in the repository and of the policy catalogue.

Together, the three components of the CTVE 101 Toolkit will offer smaller platforms a technical, policy, and enforcement suite to support with their CTVE operations. CCF will be working to build out the CTVE 101 Toolkit over the coming months, and will continue to engage with tech companies, trust and safety experts, and CTVE experts to collect feedback and refine our approach.

info@christchurchcall.org

christchurchcall.org